

Before we start today...

- Create a Github account
- Share your account name with me at tomoyas@mit.edu
- If possible (or necessary), install Git in your computer (check <https://happygitwithr.com/install-git.html>)

Git for Social Scientists: Introduction to Version Control with Git

Tomoya Sasaki¹

Massachusetts Institute of Technology

November 4th, 2022

¹This slide deck is heavily inspired by the workshop materials by Suyeol Yun and Shiro Kuriwaki.

Introduction



Companies using Git

[Airbnb](#)

[Allegro](#)

[Amazon Web Services - Labs](#)

[Amazon Web Services](#)

[Amazon](#)

[Ambientia Group Oy](#)

[Apple Inc.](#)

[arkency](#)

[AT&T Open Source](#)

Companies using Github

Introduction

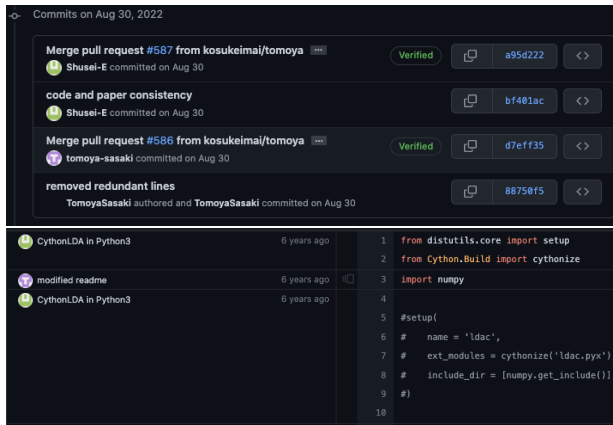
- Why Git? Why Github? Why version control?
- These are essential tools for programmers
- How about social scientists?
- My opinion: Social scientists also benefit from version control with Git
 - Increase in collaborative projects
 - Demand for clean replication materials
 - Complex data manipulation/preprocessing/analysis
- “Code and Data for the Social Sciences: A Practitioner’s Guide” by Gentzkow and Shapiro has a chapter dedicated for version control
- This workshop
 - Introduction to version control
 - Pros and cons of Git/Github
 - Brief introduction to these tools (how Git/Github works and how to use them)

What is version control?

- Version control: tracking and managing changes to file content
- Git: (the most popular) software for version control
- Github: service to host your git on the Internet (alternatives include GitLab, Bitbucket ...)
- Repository: unit of a version control project, your project folder with a subfolder named `.git`
 - Often simply called “repo”
 - `.git` folder in a repository tracks and stores every single change you make in the corresponding repository
- I focus on Git/Github because they are extremely popular than their alternatives

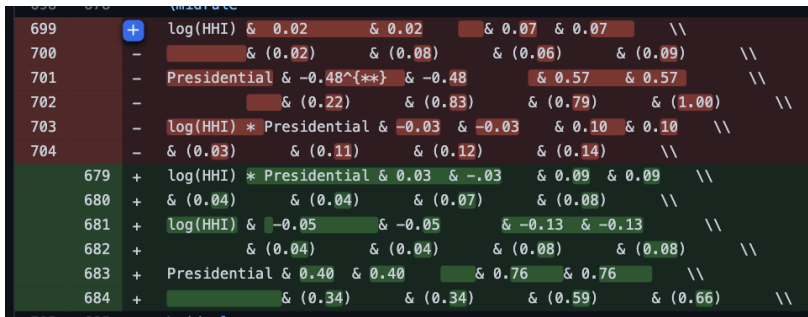
Benefit of Git/Github: Tracking who/how/when

- You can identify
 - who made changes
 - how they made the changes
 - when they made the changes
- You can check the entire history since you created a repository and move back to previous versions easily (undo changes)
- Github can visualize them nicely
- Useful when you
 - want to revert your (particular) changes
 - work on a collaborative project
- You don't need to keep
 - different versions of the same file:
clean_data_1104.R, clean_data_1020.R
 - the same file edited by different people:
clean_data_tomoya.R, clean_data_adam.R



Benefit of Git/Github: Tracking who/how/when

- You can check how the results change when we try different specification
- Easy to track which part of the results changed

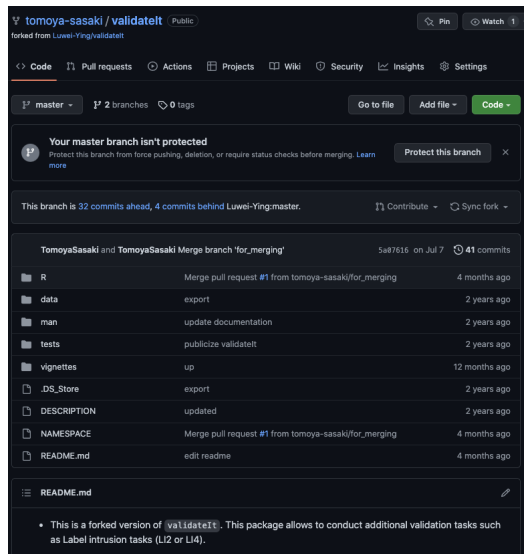


The screenshot shows a regression results table with Git diff highlights. The table has columns for coefficients, standard errors in parentheses, and t-statistics in brackets. The rows are numbered 699 to 704 and 679 to 684. The highlights indicate changes between versions: red for deletions and green for additions.

Line	Sign	Variable	Coefficient	SE	t-stat	Other
699	+	log(HHI)	0.02	0.02	0.07	0.07
700	-		(0.02)	(0.08)	(0.06)	(0.09)
701	-	Presidential	-0.48^{**}	-0.48	0.57	0.57
702	-		(0.22)	(0.83)	(0.79)	(1.00)
703	-	log(HHI) * Presidential	-0.03	-0.03	0.10	0.10
704	-		(0.03)	(0.11)	(0.12)	(0.14)
679	+	log(HHI) * Presidential	0.03	-0.03	0.09	0.09
680	+		(0.04)	(0.04)	(0.07)	(0.08)
681	+	log(HHI)	-0.05	-0.05	-0.13	-0.13
682	+		(0.04)	(0.04)	(0.08)	(0.08)
683	+	Presidential	0.40	0.40	0.76	0.76
684	+		(0.34)	(0.34)	(0.59)	(0.66)

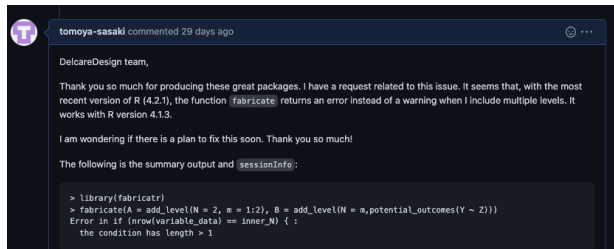
Other side benefits of Git/Github

- Hosting a customizable website (free, no ads, tons of templates)
- Contribute to software packages hosting on Github
- Tweak a package developed by someone else for your own purposes
- Send a request to package developer (often happens at “Issue”)
- Nice integration with popular apps/websites such as RStudio and Overleaf



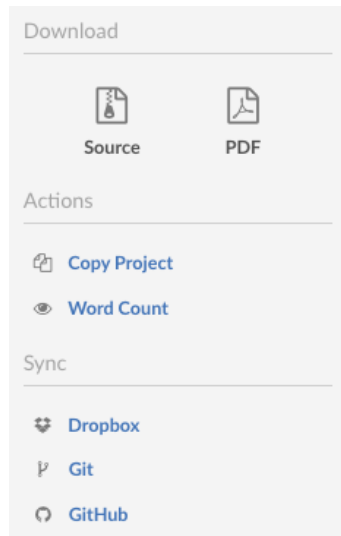
Other side benefits of Git/Github

- Hosting a customizable website (free, no ads, tons of templates)
- Contribute to software packages hosting on Github
- Tweak a package developed by someone else for your own purposes
- Send a request to package developer (often happens at “Issue”)
- Nice integration with popular apps/websites such as RStudio and Overleaf



Other side benefits of Git/Github

- Hosting a customizable website (free, no ads, tons of templates)
- Contribute to software packages hosting on Github
- Tweak a package developed by someone else for your own purposes
- Send a request to package developer (often happens at “Issue”)
- Nice integration with popular apps/websites such as RStudio and Overleaf



Limitations: not suitable for tracking large files

- Github imposes file size limit:
 - 25 MB per file limit (you can change this limit up to 100MB by changing setup)
 - 1GB per repository limit
- Remember that `.git` tracks and stores all the change you make in a repository
~> if you store a huge file in the repository and let `.git` tracks its changes, the `.git` folder can grow quite huge
- Use `.gitignore` to specify files that Git should not track “ignore”

```
.gitignore  
*.csv # ignore csv files  
/data/ # ignore any file in data folder
```
- Include huge files as well as sensitive files that contain password, API key etc in `.gitignore`


Limitations: not great if you want to track non-text files

- Git cannot track line by line changes for non-text files such as PDF, Microsoft Word/Excel/Powerpoint, JPG, ...
- Note that Git still tracks changes
- The value of Git/Github is limited
- In the right example, Git/Github recognizes the changes as the changes in file sizes
→ even though you update a figure in PDF or PNG format, Git/Github might not recognize it unless the file size changes...



How Git tracks files

- Git tracks changes in files with “**commit**”
- Changes between commits are called “**diff**”
- Each commit is a snapshot of your repository at specific times
- Commit has a human-readable message and an (first few characters of) commit ID
- Commit is not automatic and you actively make a “commit” with a message
- Ideally commit should be atomic units of change that represent a specific idea
- You need to “**stage**” the files you want to add to commit (more later)




- 8bacc32: “test commit”
- e9f128a: “edit analysis code”
- 0d492rg: “add new analysis”

Github and more on repos

- If you just work on your local repos (and only on one computer), you don't need Github
- Repos: unit of a version control project, contains a folder with a subfolder named `.git`
 - Local repos: repos (folder) in your own computer, ordinary project folder with subfolder named `.git`
 - Remote repos: repos (folder) in a web hosting services such as Github
- Remote repos such as Github...
 - come with unique URL (`https://github.com/username/reponame.git`)
 - are basically a copy of your local repos
 - nicely visualize your commits and also serve as a backup of your codes (remember, not for data)
 - useful when you work on the same project with multiple computers
- You can choose remote repo to be public or private
- Github accounts let you create unlimited number of private repos with upto 3 collaborators
unlimited number of public repos with unlimited number of collaborators

How to interact with remote repos on Github

- Interaction with remote repos: **pull** and **push**
- You **push** the commits you made to remote repos
- You **pull** new commits on remote repos
- Technically “pull” does **fetch** and **merge**



How to use Git/Github

- So far, we have covered some fundamental concepts around Git/Github
- But how do we initialize a repo and commit/pull/push in practice?
- Command line, Github desktop app etc
- This workshop focuses on RStudio
- I personally use command line

General (suggestive) initialization

- ❶ Establish a connection between your computer and Github
 - ❷ Create a repo on Github (create a public one for now)
 - ❸ **Clone** your remote repo on Github to your computer (can be any location). This is your local repo.
 - ❹ Edit `.gitignore` to make sure that Git won't track any sensitive or huge file
- If you have any existing project that you want to track with Git, create a remote repo on Github, clone it to your computer (i.e., local repo), and move codes and folders into the local repo

General (suggestive) daily workflow


- 1 Edit, add, or delete files in your repo
- 2 Stage files you want to add to commit and make a commit when you make some changes that represent a specific idea or solve particular issues
 - Add lines to clean data
 - Add lines to visualize regression results
 - Delete lines that are irrelevant anymore
- 3 Push when you make a set of commits or when you call it a day
- 4 Pull any new commits in the remote repo if there are commits and pushes to the remote repo from other computers
- 5 Repeat the process above
 - Push at least once a day if you make any edit
 - Let's try out with Github and RStudio

Setup authentication

- When we interact with a remote Git server, such as GitHub, we need to provide credentials
- Github provides two authentication methods, HTTPS and SSH
- I recommend HTTPS and will use HTTPS in this workshop
- Procedure
 - ➊ Go to <https://github.com/settings/tokens> and click “Generate token”
 - ➋ Decide the scope. Choose (at least) “repo”, “user”, and “workflow”.
 - ➌ Click “Generate token”. This is the same as the password to interact with Git. You shouldn’t show this to anyone.
 - ➍ Copy the generated strings (PAT, personal access token)
 - ➎ Open RStudio and run `gitcreds::gitcreds_set()` (Install the `gitcreds` package beforehand)
 - ➏ Paste your PAT
 - ➐ RStudio stores and remembers PAT for you

How Git/Github handle collaborative projects

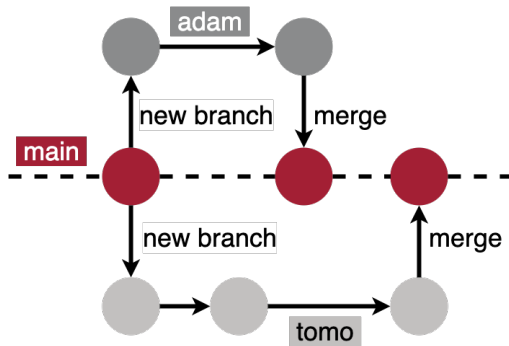
- Easy collaboration is a distinct feature of Git
- Key concepts: **branch** and **fork**
- Branches are parallel universe of a repository
- Main branch is the default branch and should contain stable version
- You contribute to a repo by creating a new branch in isolation from changes that other people are making to the repo
- Whenever you work on a collaborative project, you should always create a new branch



- Each circle represents “commit”
- In this example, Adam and Tomo work on the same project and each of them creates a branch

How Git/Github handle collaborative projects

- Once your work is done, open **pull request**, and **merge** your branch into the default branch
- Pull request is an opportunity to discuss or check (**review**) if the changes you make in your branch won't break the default branch
- By merging your branch into the default branch, any edit you make in your branch will be reflected in the default branch
- When both of them finish their work, they open a pull request and merge their branch into the main branch



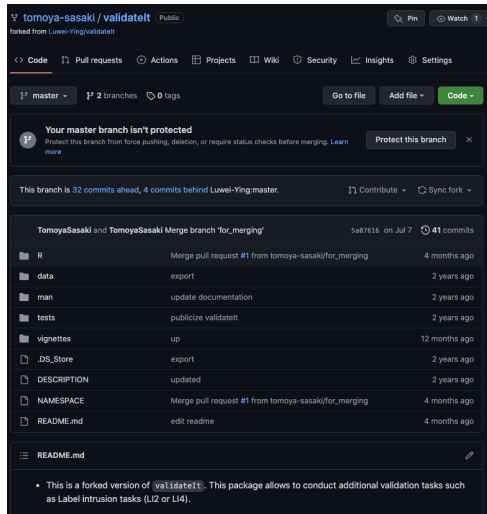
- Each circle represent “commit”
- In this example, Adam and Tomo work on the same project and each of them creates a branch

General (suggestive) daily workflow with multiple branches

- ❶ Create a new branch (usually from the default branch)
- ❷ In your branch...
 - ❶ Edit, add, or delete files in your repo
 - ❷ Stage files you want to add to commit and make a commit when you make some changes that represent a specific idea or solve particular issues
 - ❸ Push when you make a set of commits or when you call it a day
 - ❹ Repeat the process above
- ❸ Create a pull request when you finish your project
- ❹ Merge your branch into the default branch
 - Let's try out with Github and RStudio

Working on someone else's repo

- Even if you have a public repository on Github, only those who have permission can directory push and merge to your repo
- What to do if you want to contribute to someone else's repo or borrow their idea and tailor for your own project?
- **Fork** their project to your Github account
 - Original: <https://github.com/xxx/reponame.git>
 - Forked repo:
<https://github.com/yourusername/reponame.git>



Comparison and relationship to alternatives (Dropbox, Google Drive)

- Unlike Dropbox or Google Drive, Git/Github does not automatically track changes
 ~> you need to take (small) action (mostly) when you
- However, you can make your project folder clean with commits and branches
- You can revert your changes easily and flexibly if you make frequent commits
 ~> Dropbox and Google Drive only preserve once-in-a-day snapshots
- Putting Git local repos in Dropbox or Google Drive can be very tricky when your Dropbox/Google Drive account is synced to multiple computers
- Git does not like autosync (autoupdate) by Dropbox/Google Drive
- But you probably want to rely on these other services to store huge data that Git does not track
- In my setup, I put Git repos in Dropbox but sync them only with my main computer and the same Git repos in my other computers are located in non-Dropbox folders

What's next?

- Working with private repos
- (Gradually) Learning how to use git in command lines
- Probably you mess up once in a while. I still make mistakes sometimes
- You can fix them easily by reverting

Useful resources

- Happy Git and GitHub for the useR: <https://happygitwithr.com/index.html>
- Git Guide: <https://github.com/git-guides>
- git-vs-dropbox: <https://michaelstepner.com/blog/git-vs-dropbox/>